

PaperTunedLLM Project Documentation

1. Executive Summary

PaperTunedLLM is a production-ready AI Research Assistant specialized in Machine Learning research papers. It is built using modern open-source Large Language Model (LLM) technologies and demonstrates the complete lifecycle of adapting an LLM, encompassing fine-tuning, evaluation, quantization, retrieval-augmented generation (RAG), and deployment.

2. Architecture and Core Components

The project architecture is composed of several functional layers integrated to provide a robust research assistant.

2.1 Data & Base Model

- Base LLM:** Qwen 2.5 3B, providing the foundation for general understanding.
- Dataset:** Existing public scientific datasets (e.g., QASPER) utilized for supervised fine-tuning.

2.2 Training Pipeline

- Methodology:** QLoRA for efficient supervised fine-tuning, yielding the specialized PaperTunedLLM.
- Frameworks:** LlamaFactory and Unsloth for optimized training.

2.3 Evaluation & Optimization

- Evaluation:** Rigorous comparative analysis against the base model, measuring scientific QA performance, paper understanding, hallucination rates, and reasoning quality.
- Quantization:** AWQ quantization to reduce GPU memory footprint while maintaining accuracy, ensuring an efficient production model.

2.4 Retrieval-Augmented Generation (RAG)

- Embedding Model:** BGE-M3 for converting document chunks into vector representations.
- Vector Database:** Qdrant for storing embeddings and enabling rapid semantic retrieval.

2.5 Application Layer

- Inference Engine:** Text Generation Inference (TGI) for serving the quantized model efficiently in a production setting.
- Backend:** FastAPI to handle requests, orchestrate RAG retrieval, and interface with the LLM.
- Frontend:** React and Next.js providing the user interface for uploading papers, querying the knowledge base, and analyzing research.

2.6 Deployment

- Infrastructure:** Docker Compose for a unified, single-command deployment workflow.

3. System Architecture Diagram

```
%%{init: {'theme': 'base', 'themeVariables': { 'primaryColor': '#ffffff', 'primaryTextColor': '#000000',  
'primaryBorderColor': '#000000', 'lineColor': '#000000', 'secondaryColor': '#ffffff', 'tertiaryColor':  
      '#ffffff' }}}%%  
flowchart TD  
    %% Styling  
    classDef bw fill:#fff,stroke:#000,stroke-width:2px,color:#000;  
    classDef db fill:#fff,stroke:#000,stroke-width:2px,color:#000,shape:cylinder;  
  
    %% Nodes
```

```

User[User / Researcher]:::bw
UI[Frontend: React / Next.js]:::bw
API[Backend: FastAPI]:::bw

subgraph RAG_Pipeline [RAG Pipeline]
  Embed[Embedding: BGE-M3]:::bw
  VDB[(Vector DB: Qdrant)]:::db
end

subgraph LLM_Inference [LLM & Inference]
  TGI[Inference: TGI Engine]:::bw
  Model[PaperTunedLLM AWQ Quantized]:::bw
end

%% Connections
User <-->|Queries & Uploads| UI
UI <-->|REST API| API

API -->|1. Document Chunks| Embed
Embed -->|2. Vector Embeddings| VDB
API -->|3. Query Embedding| Embed
VDB -->|4. Retrieve Context| API

API -->|5. Prompt + Context| TGI
TGI <--> Model
TGI -->|6. Generated Response| API

```

4. System Workflows

4.1 Training Flow

1. **Public Dataset** processing.
2. **QLoRA Fine-Tuning** using LlamaFactory/Unsloth.
3. **Evaluation** of model metrics against baseline.
4. **AWQ Quantization** for optimization.
5. **Production Model** readiness.

4.2 Retrieval (RAG) Flow

1. **Scientific Papers** ingestion.
2. **Chunking** the documents into logical pieces.
3. **BGE-M3 Embedding** generation.
4. **Qdrant Vector DB** storage and indexing.
5. **Retriever** fetching context.
6. **PaperTunedLLM** generation based on retrieved context.

4.3 User Interaction Flow

1. Researcher interacts with the Web UI.
2. Request is forwarded to the FastAPI backend.
3. BGE-M3 generates a query embedding.
4. Qdrant performs a semantic search to identify relevant document chunks.
5. PaperTunedLLM receives the prompt augmented with the retrieved chunks.
6. A grounded, accurate response is generated and returned to the user.

5. Technology Stack Overview

Component	Technology
Base Model	Qwen 2.5 3B
Fine-Tuning	QLoRA
Quantization	AWQ
Embedding Model	BGE-M3
Vector Database	Qdrant
Inference	Text Generation Inference (TGI)
Backend Framework	FastAPI
Frontend Framework	React / Next.js
Containerization	Docker